



reSiliant coMputer archItectures
and LIfe Sciences



Politecnico
di Torino

Department of Control and
Computer Engineering



PROCESS MANAGEMENT

STEFANO DI CARLO

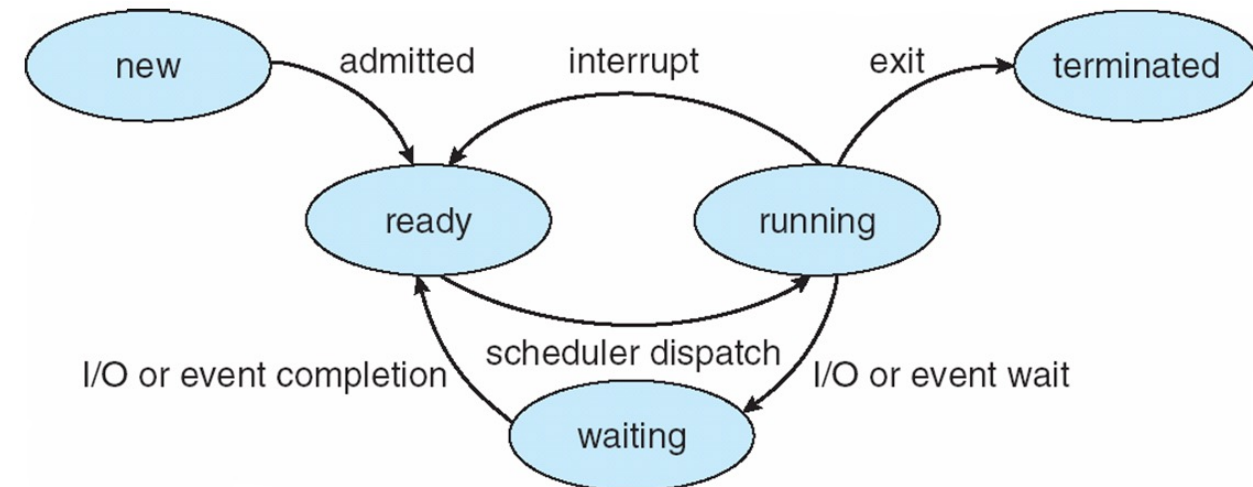
PROCESS DEFINITION

- ▶ **Process** – a program in execution
 - ▶ process execution must progress in sequential fashion.
 - ▶ No parallel execution of instructions of a single process
- ▶ **Multiple parts**
 - ▶ The program code, also called **text section**
 - ▶ Current activity including **program counter**, processor registers
 - ▶ **Stack** containing temporary data
 - ▶ Function parameters, return addresses, local variables
 - ▶ **Data section** containing global variables
 - ▶ **Heap** containing memory dynamically allocated during run time

PROCESS STATE

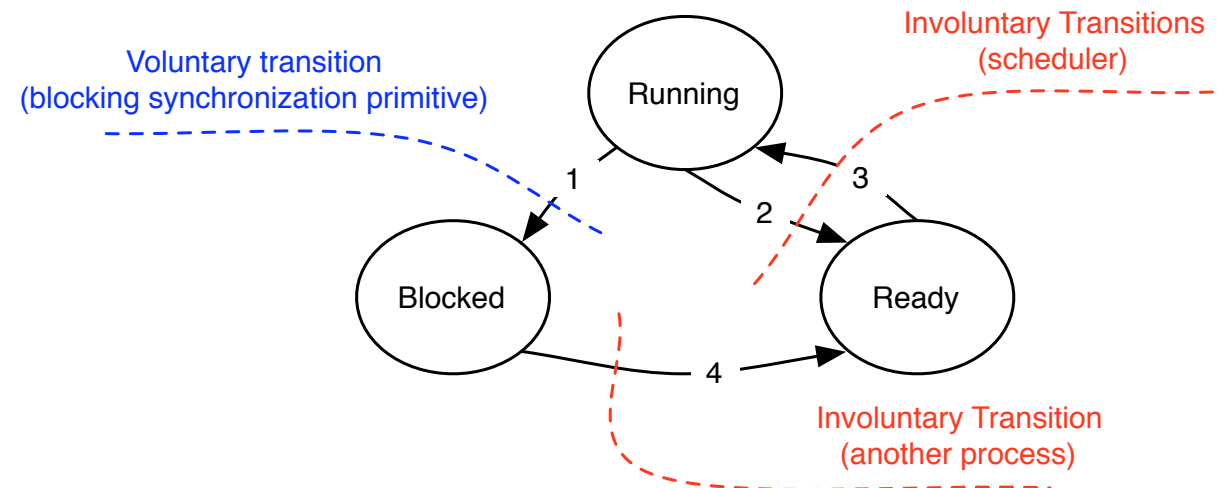
- ▶ As a process executes, it changes state
 - ▶ New: The process is being created
 - ▶ Running: Instructions are being executed
 - ▶ Waiting: The process is waiting for some event to occur
 - ▶ Ready: The process is waiting to be assigned to a processor
 - ▶ Terminated: The process has finished execution

- ▶ Process State Diagram (PSD)



PSD TRANSITIONS

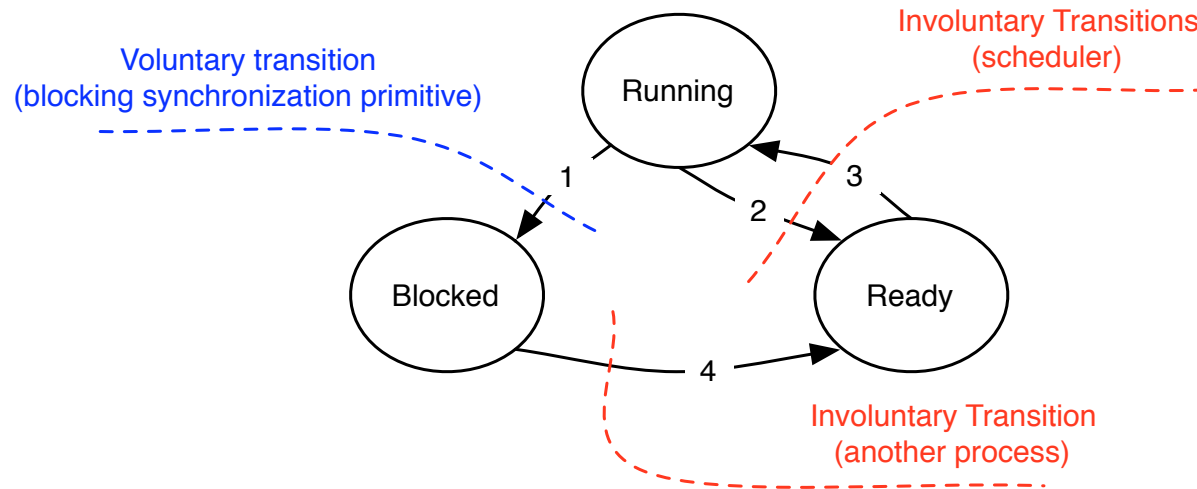
- ▶ Given N processes, and one processor, at any given time:
 - ▶ 1 process is in the running state
 - ▶ M processes can be blocked waiting for a resource to become available to resume the execution
 - ▶ $N-M-1$ are ready to be executed waiting to access the processor
- ▶ Transition 1 occurs when a process discovers that it cannot continue
 - ▶ For example it needs to use a portion of shared memory a now-ready process reserved for its own use



PSD TRANSITIONS

- ▶ Transition 4 occurs when the event the process was waiting for occurs
 - ▶ For example, the running process sets free the shared memory it previously locked, which the blocked process was waiting for

- ▶ Transitions 2 and 3 are caused by the operating system
 - ▶ Preemptive scheduler: process is moved from running to ready after a certain time quantum (time slice) is expired
 - ▶ Cooperative scheduler: process voluntary moves from running to ready

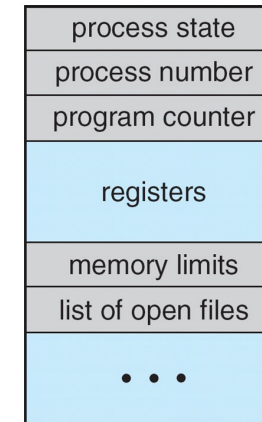


PROCESS CONTROL BLOCK (PCB)

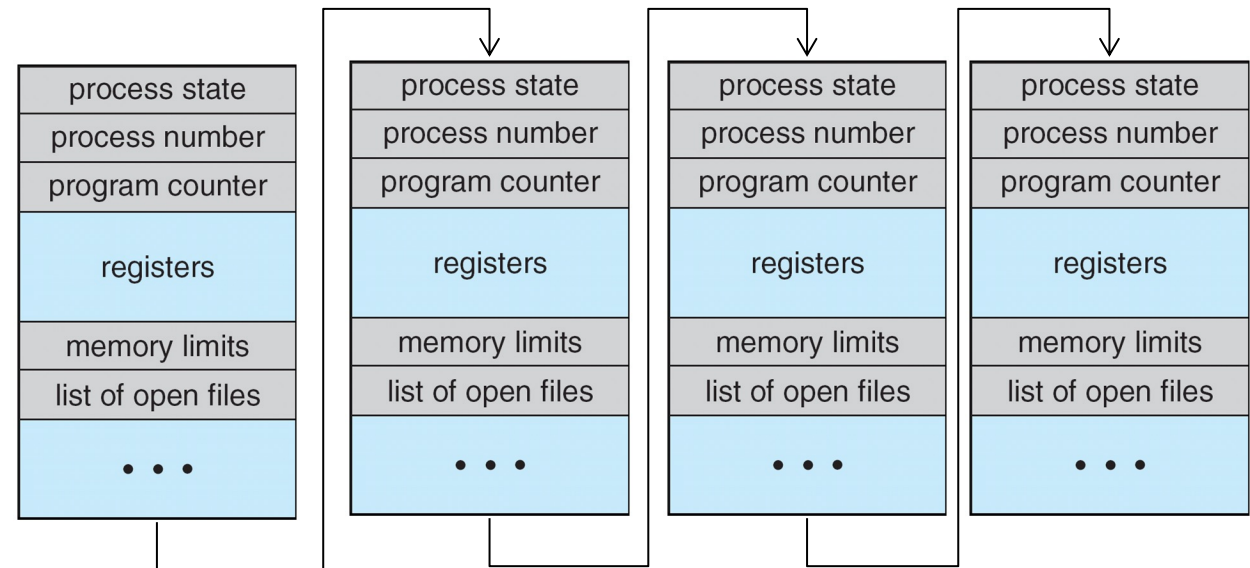
- ▶ The CPU Manager describes each process using the PCB containing:

- ▶ Process state
- ▶ Program counter
- ▶ CPU registers
- ▶ CPU scheduling information
- ▶ Memory-management information
- ▶ Accounting information
- ▶ I/O status information

- ▶ Example of PCB



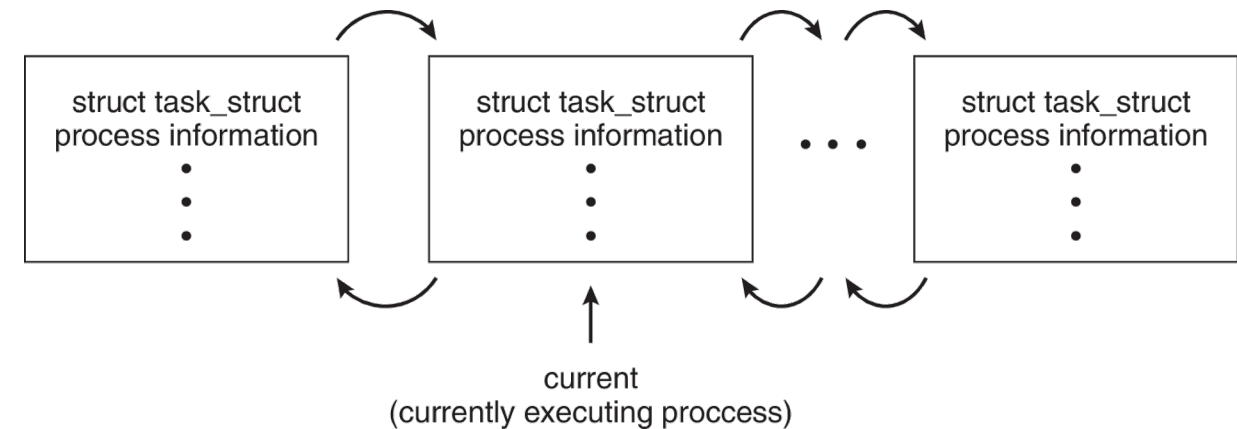
- ▶ One PCB is maintained for each process



PROCESS REPRESENTATION IN LINUX


Represented by the C structure

```
struct task_struct {
    pid_t pid;          /* process identifier */
    long state;         /* state of the process */
    unsigned int time_slice /* scheduling
information */
    struct task_struct *parent; /* this
process's parent */
    struct list_head children; /* this
process's children */
    struct files_struct *files; /* list of open
files */
    struct mm_struct *mm;    /* address space
of this process */
}
```



THE CPU SCHEDULER

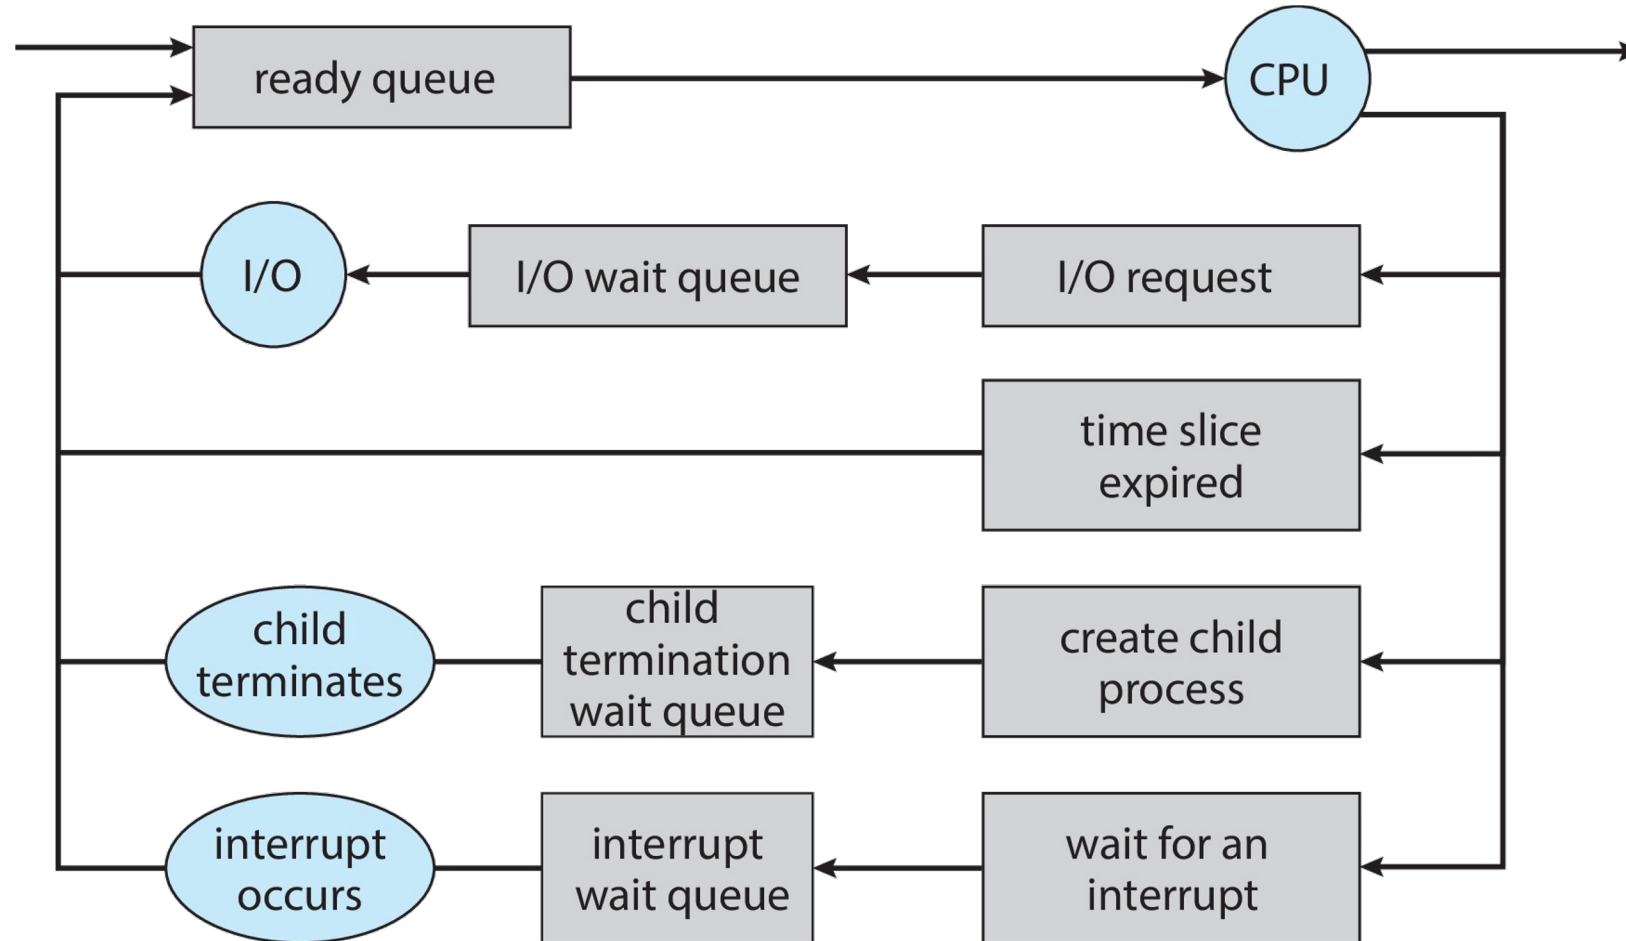
- ▶ It is the part of the CPU manager that implements the process state transitions and decides which process must run
- ▶ The goal is to maximize CPU use
- ▶ Maintains scheduling queues of processes
 - ▶ Ready queue – set of all processes residing in main memory, ready and waiting to execute
 - ▶ Wait queues – set of processes waiting for an event (i.e., I/O)
 - ▶ Processes migrate among the various queues



SCHEDULING CRITERIA

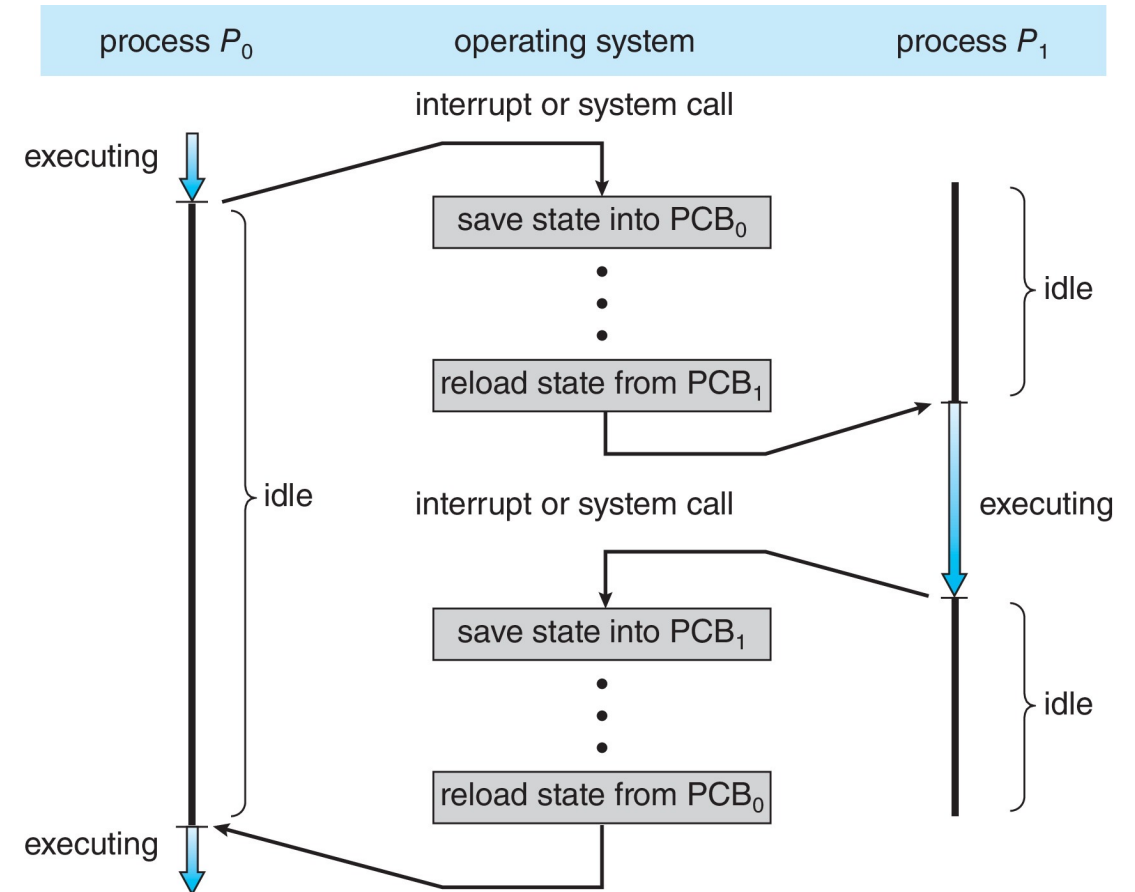
- ▶ CPU utilization – keep the CPU as busy as possible
- ▶ Throughput – # of processes that complete their execution per time unit
- ▶ Turnaround time – amount of time to execute a particular process
- ▶ Waiting time – amount of time a process has been waiting in the ready queue
- ▶ Response time – amount of time it takes from when a request was submitted until the first response is produced.

REPRESENTATION OF PROCESS SCHEDULING



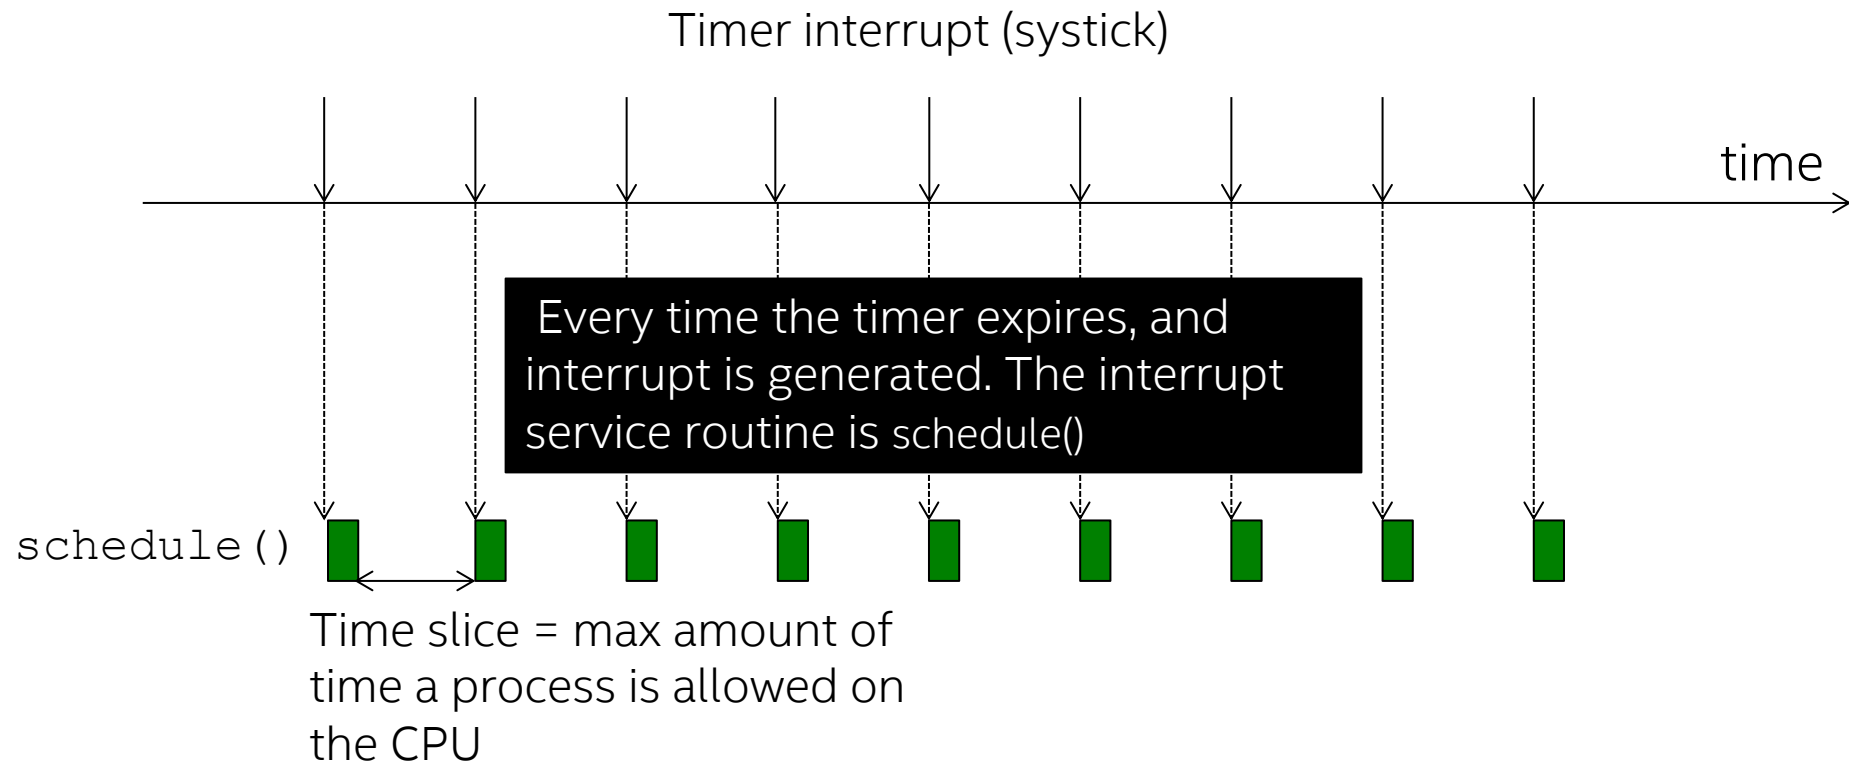
CONTEXT SWITCH

- ▶ When CPU switches to another process, the system must **save the state** of the old process and load the **saved state** for the new process via a **context switch**
- ▶ **Context** of a process represented in the PCB
- ▶ Context-switch time is pure overhead; the system does no useful work while switching
 - ▶ The more complex the OS and the PCB → the longer the context switch
- ▶ Time dependent on hardware support
 - ▶ Some hardware provides multiple sets of registers per CPU → multiple contexts loaded at once




SCHEDULER ACTIVATION

- It is called periodically (via timer interrupt) or in response to events




PREEMPTION

- ▶ Operation performed to evict a running process from the CPU
- ▶ Example: upon reaching the end of the time slice, process A is moved from running to ready, while process B is moved from ready to running



BLOCKING

- ▶ Preemption may take place as a result of an operation the running task performs
- ▶ Example: Process A is moved to blocked state, while process B is set running



Process A executed an instruction the blocks it

WHAT DOES A PROCESS (OR TASK) LOOK LIKE?

- ▶ Basic task
 - ▶ Sequence of statements executed once for each instance of the task (e.g., `do_instance()`)
 - ▶ Instance of the task = execution of the function `x()`
 - ▶ The task starts with the first instruction of function `x()`
 - ▶ The task terminates after the last statement of function `x()`
 - ▶ An initialization function `x_init()` is executed once for setting up the memory used by task `x`
 - ▶ If task `x` needs to keep in memory data to be used by different instances, global variables shall be used (e.g., `persistent_data`)

```
int persistent_data;
```

```
Task x()
```

```
{  
    do_instance();  
}
```

```
x_init()
```

```
{  
    initialization();  
    ...  
}
```

WHAT DOES A PROCESS (OR TASK) LOOK LIKE?

- ▶ Extended task
 - ▶ It is a function that starts once, and never ends
 - ▶ Data can be local variables
 - ▶ Initialization operations are performed once, before starting the end-less loop
 - ▶ The end-less loop implements the operation of the task (e.g., `do_instance()`)
 - ▶ The end-less loop typically contains a statement to block the task until it is needed (e.g., `WaitEvent()`)
 - ▶ A certain amount of time is elapsed
 - ▶ A resource becomes available
 - ▶ ...
 - ▶ When the task is blocked others may run



```
Task x()
{
    int local;

    initialization();

    for (;;) /* End-less loop */
    {
        WaitEvent();
        do_instance();
    }
}
```


SCHEDULING ALGORITHMS

- ▶ It is the criteria to pick up a process among those ready to make it running
- ▶ Simplest form of scheduling: Round Robin
 - ▶ Processes are inserted into a FIFO queue {B, F, D, G, A}
 - ▶ The top of the queue is executed (see figure a)
 - ▶ When exiting from running state it is queued to the last position of the queue (see figure b)



ROUND ROBIN SCHEDULER

► Possible implementation

```
Schedule(ReadyList R, RunningTask T)  
{  
    T->TCB.state = READY;  
    save_context(T->TCB);  
  
    append_to_list( T, R );  
  
    Q = top_of_list( R );  
    restore_context(Q->TCB);  
    Q->TCB.state = RUNNING;  
}
```

Running
Task



List of Ready
Tasks R



Top of the
list




ROUND ROBIN SCHEDULER

► Possible implementation

```
Schedule(ReadyList R, RunningTask T)
{
    T->TCB.state = READY;
    save_context(T->TCB);

    append_to_list( T, R );

    Q = top_of_list( R );
    restore_context(Q->TCB);
    Q->TCB.state = RUNNING;
}
```




Running
Task

List of Ready
Tasks R

T_0

T_1 T_2 T_3 T_0

Top of the
list



ROUND ROBIN SCHEDULER

► Possible implementation


```
Schedule(ReadyList R, RunningTask T)
{
    T->TCB.state = READY;
    save_context(T->TCB);

    append_to_list( T, R );

    Q = top_of_list( R );
    restore_context(Q->TCB);
    Q->TCB.state = RUNNING;
}
```



Running
Task

List of Ready
Tasks R



SCHEDULING ALGORITHMS

- ▶ In Round Robin all processes are equal



- ▶ What if B has a more important task to do with respect to the others?
 - ▶ After it run, it has to wait 4 time slices before running again
- ▶ Solution:
 - ▶ To differentiate processes assigning a weight factor → priority
 - ▶ To adopt a scheduling approach based on priority → priority-based scheduling

PRIORITY-BASED SCHEDULER

```
Schedule(ReadyList R, RunningTask T)
{
    Q = top_of_list( R );

    if( Priority(Q) > Priority(T) )
    {
        T->TCB.state = READY;
        save_context(T->TCB);

        append_to_list( T, R );

        restore_context(Q->TCB);
        Q->TCB.state = RUNNING;
    }
}
```


- ▶ Preemption of a running task & context switch happen only if a ready task exists whose priority is greater than that of the running task
- ▶ If the running task has priority equal to that of the highest priority ready task preemption does not happen
- ▶ Highest priority running task is preempted when it has to be blocked

SCHEDULING ALGORITHMS

- ▶ A priority function is defined which returns numerical value T for process p :
 - ▶ $T = \text{Priority}(p)$
- ▶ Static priority: unchanged for lifetime of p
- ▶ Dynamic priority: changes at runtime

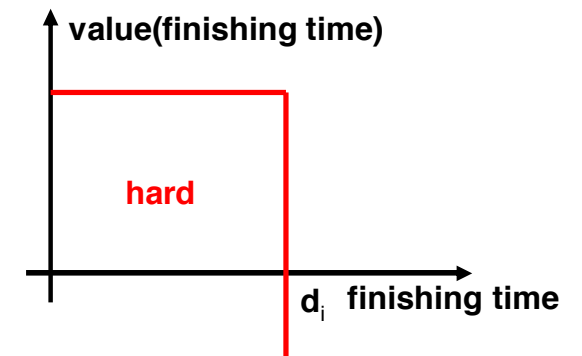



TYPE OF SYSTEMS

- ▶ Non real-time system: computer that has to respond to external events correctly
 - ▶ Example: given $x(t)$ at time $t \rightarrow$ output y must be delivered anytime
- ▶ Real-time system: computer that has to respond to external events both correctly and within a finite, specified period of time called deadline
 - ▶ Example: given $x(t)$ at time $t \rightarrow$ output y must be delivered no later than $t+d$
 - ▶ Right result too late is as bad as giving wrong or no result




REAL TIME FLAVORS

- ▶ A process is hard real time if missing its deadline may cause catastrophic consequences on the environment under control
- ▶ A process is firm real time if missing its deadline makes the result useless, but missing does not cause serious damage
- ▶ A process is soft real time if meeting its deadline is desirable (e.g. for performance reasons) but missing does not cause serious damage




REAL TIME VS FAST

- ▶ The objective of a real-time system is to guarantee the timing behavior of each individual process



- ▶ The objective of a fast system is to minimize the average time of a set of processes takes to complete



QUESTIONS?

THANK YOU!

