



reSilient coMputer archItectures
and LiFE Sciences



Politecnico
di Torino

Department of Control and
Computer Engineering



SCHEDULING PREDICTABILITY

STEFANO DI CARLO

PREDICTABLE HARDWARE SUPPORT

- ▶ Predictability: the OS should guarantee that all critical timing constraints are met
- ▶ Predictability requires determinism
 - ▶ Deterministic algorithm is an algorithm which, given a particular input, will always produce the same output, with the underlying machine always passing through the same sequence of states
- ▶ Determinism is influenced by
 - ▶ The architectural features of the hardware
 - ▶ The scheduling algorithm
 - ▶ The IPC mechanisms
 - ▶ ...

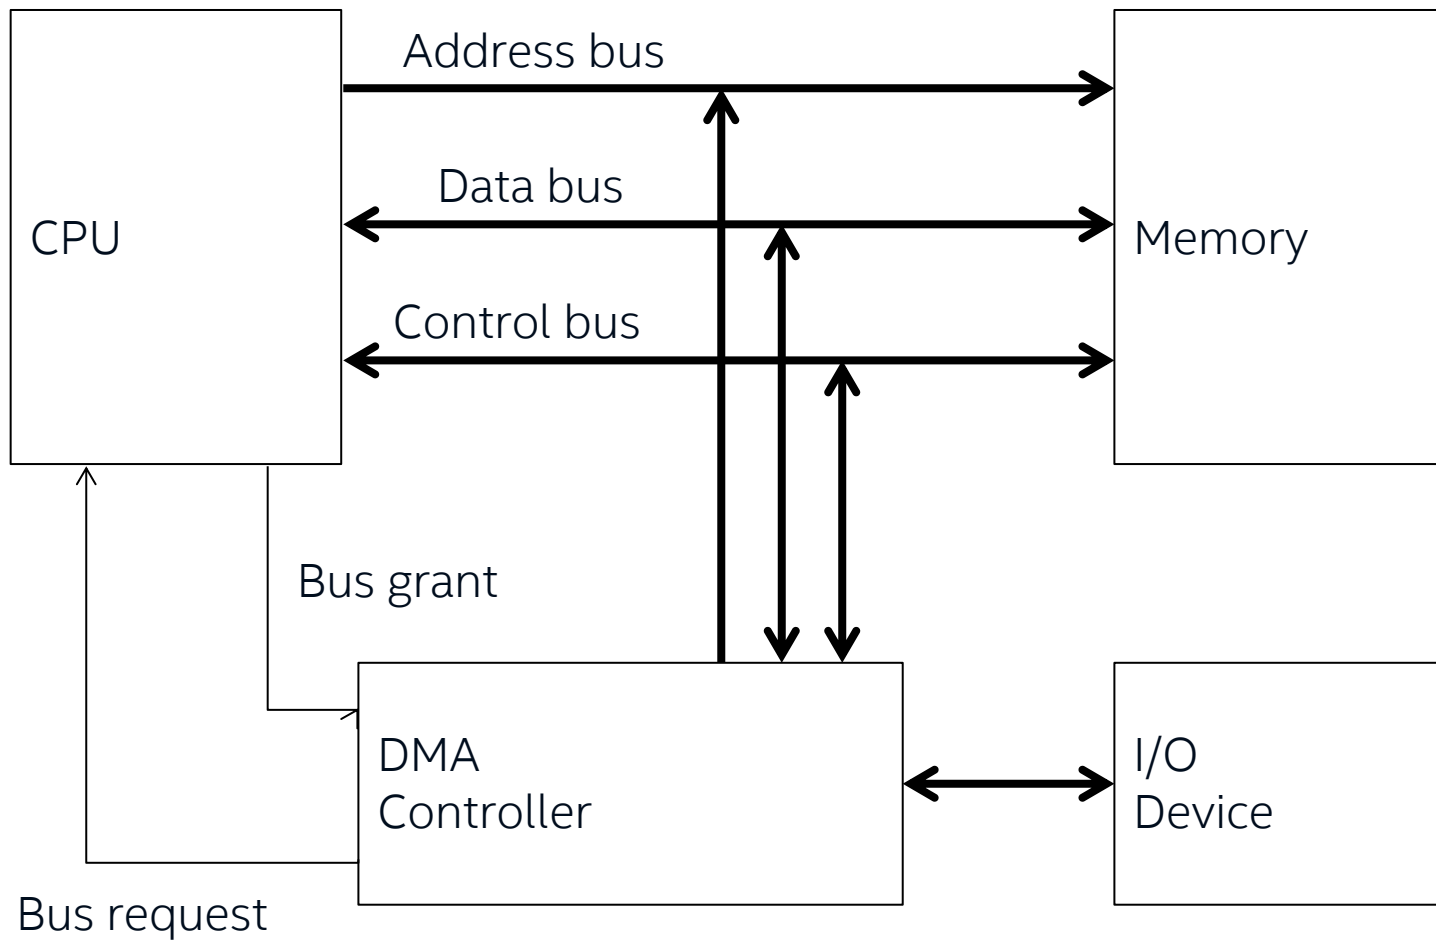
HW/SW SOURCES OF NONDETERMINISM

- ▶ Direct memory access (DMA)
- ▶ Cache memory
- ▶ Interrupts
- ▶ System calls
- ▶ Semaphores
- ▶ Memory management
- ▶ Programming language

DIRECT MEMORY ACCESS (DMA)

- ▶ DMA is used by many peripheral devices to transfer data between the device and the main memory
 - ▶ It relieves the CPU of the task of controlling the I/O transfer
- ▶ The CPU and the I/O device share the same bus
- ▶ The CPU is blocked when the DMA device is performing a data transfer

DIRECT MEMORY ACCESS (DMA)



DIRECT MEMORY ACCESS (DMA)

- ▶ Data transfer mode
 - ▶ Burst
 - ▶ An entire block of data is transferred in one contiguous sequence
 - ▶ The CPU remains inactive for relatively long periods of time (until the whole transfer is completed)
 - ▶ Cycle stealing
 - ▶ DMA transfers one byte of data and then release the bus returning control to the CPU
 - ▶ It continually issues requests, transferring one byte of data per request, until it has transferred the entire block of data
 - ▶ It takes longer to transfer data/the CPU is blocked for less time
- ▶ Transparent
 - ▶ DMA transfers data when the CPU is performing operations that do not use the system buses

DIRECT MEMORY ACCESS (DMA)

- ▶ DMA makes difficult to predict task response time
- ▶ The time the task remains blocked depends on
 - ▶ The size of the block of data
 - ▶ The transfer mode
- ▶ Real-time oriented transfer mode
 - ▶ Time-sliced method
 - ▶ Each memory cycle is split into two adjacent time slots: one reserved for the CPU and the other for the DMA device
 - ▶ More time expensive but more predictable
 - ▶ CPU and DMA never interfere with each other
 - ▶ Task completion time does not depend on the DMA transfer

CACHE MEMORY

- ▶ Fast memory inserted as a buffer between the CPU and the RAM to speed up task execution
- ▶ At each memory access the hardware checks whether the requested information is stored in the cache:
 - ▶ Cache hit → data are read from the cache (very fast)
 - ▶ Cache miss → data are read from the RAM; the data is copied in the cache along with a set of adjacent locations (very slow)
- ▶ Exploit the program locality concept

CACHE MEMORY

- ▶ Sources of nondeterminism
 - ▶ Hit-ratio improves the performance on average
 - ▶ 80% of the time the data is in cache, 20% it is not
 - ▶ How to compute precisely the worst-case execution time?
 - ▶ In preemptive system, hit-ratio is affected by the number of preemption
 - ▶ Preemption destroys program locality, and increases the number of cache misses
- ▶ Example
 - ▶ PowerPC MPC7410 with 2 MByte two-way associative L2 cache
 - ▶ 33% WCET increment due to cache interference w.r.t. non-preemptive mode

CACHE MEMORY

- ▶ Solutions
 - ▶ Disable the cache memory
 - ▶ Impact the performance
 - ▶ More deterministic behavior
 - ▶ Overestimate the WCET
 - ▶ Does not impact the performance
 - ▶ Uses only a fraction of the available CPU time

INTERRUPTS

- ▶ They are event generated by I/O devices to get the attention of the CPU when a data is available
- ▶ The arrival of an interrupt causes the execution of an interrupt service routine
 - ▶ Fragment of code dedicated to the management of its associated I/O device
- ▶ Example: to get data from an I/O device, each task must
 - ▶ Enable the hardware to generate interrupts
 - ▶ Wait for the interrupt
 - ▶ Read the data from a memory buffer shared with the driver

INTERRUPTS

- ▶ In general-purpose OS
 - ▶ Interrupts are served using a fixed priority scheme
 - ▶ Each driver is scheduled based on a static priority, higher than process priorities
 - ▶ Interrupt service routines usually deal with I/O devices that have real-time constraints, whereas most application programs do not
- ▶ In real-time OS
 - ▶ Tasks (e.g., control process) could be more urgent than an interrupt service routine
 - ▶ It is very difficult to bound a priori the number of interrupts that a task may experience
 - ▶ The delay introduced by interrupts on tasks execution is unpredictable

INTERRUPTS

- ▶ Solution 1
 - ▶ Disable all interrupts but timer interrupt (needed for scheduling tasks)
 - ▶ I/O devices are handled using polling
 - ▶ The application task is responsible for checking I/O devices periodically for available data
 - ▶ Deterministic but requires busy waiting for communicating with I/O device
 - ▶ The I/O handling has the same priority of the application task where it is implemented

INTERRUPTS

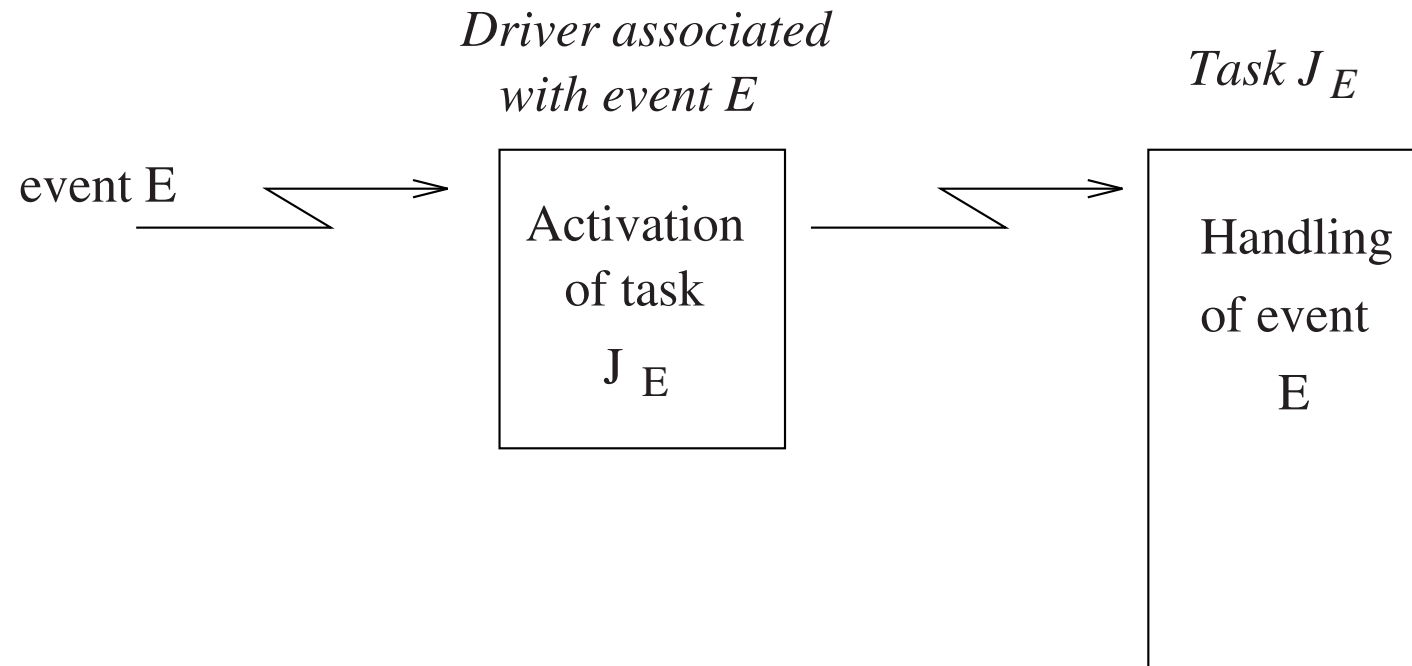
- ▶ Solution 2
 - ▶ Disable all interrupts but timer interrupt
 - ▶ Treat the I/O devices as aperiodic tasks managed through servers
 - ▶ Slow devices are multiplexed and served by a single server running at a low rate
 - ▶ Fast devices are served by dedicated servers running at higher frequencies
 - ▶ Servers are implemented by the OS
 - ▶ Still require busy waiting in the servers
 - ▶ Servers are scheduled accordingly to their priority, independently from those of application tasks

INTERRUPTS

- ▶ Solution 3
 - ▶ All interrupts are enabled
 - ▶ The interrupt handling is performed using
 - ▶ A driver: its only purpose is to activate the task associated to the I/O device
 - ▶ A task: it is responsible for handling the I/O device
 - ▶ The driver execution time is very limited → negligible impact on the tasks WCET
 - ▶ The I/O task is scheduled by the OS based on its priority, independent from those of application tasks
 - ▶ No busy waiting

INTERRUPTS

► Conceptual view of solution 3



SYSTEM CALLS

- ▶ The execution time of system calls must be characterized, and considered when evaluating if the scheduling is feasible
- ▶ System calls must be pre-emptible
 - ▶ Any non-pre-emptible section could delay the activation or the execution of critical tasks
- ▶ Example
 - ▶ Linux vanilla server is not suitable for real-time
 - ▶ The kernel functions (including system calls) is not preemptable
 - ▶ Kernel extensions are needed to make it more deterministic
 - ▶ Basically, the kernel must become pre-emptible

SEMAPHORES

- ▶ Traditional semaphores are not suitable for real-time systems → may produce the priority inversion problem
- ▶ The OS must implement into semaphores the solutions to priority inversion
- ▶ We will study this in detail later ...

MEMORY MANAGEMENT

- ▶ Memory management techniques must not introduce nondeterministic delays during the execution of real-time activities
- ▶ We will study this in detail later ...

PROGRAMMING LANGUAGE

- ▶ Restrictions that may be needed for real-time
 - ▶ No dynamic data structures: dynamic arrays must be eliminated because they would prevent a correct evaluation of the time required to allocate and deallocate memory
 - ▶ No recursion: If recursive calls were permitted, the schedulability analyzer could not determine the execution time of subprograms involving recursion or how much storage will be required during execution
 - ▶ Time-bounded loops: In order to estimate the duration of the cycles at compile time, programmer should specify for each loop construct the maximum number of iterations

QUESTIONS?

THANK YOU!

